

# The Spy In The Sandbox: Practical Cache Attacks In Javascript And Their Implication

Yossef Oren   Vasileios P. Kemerlis   Simha Sethumadhavan   Angelos D. Keromytis

Apresentação por Renan Greca

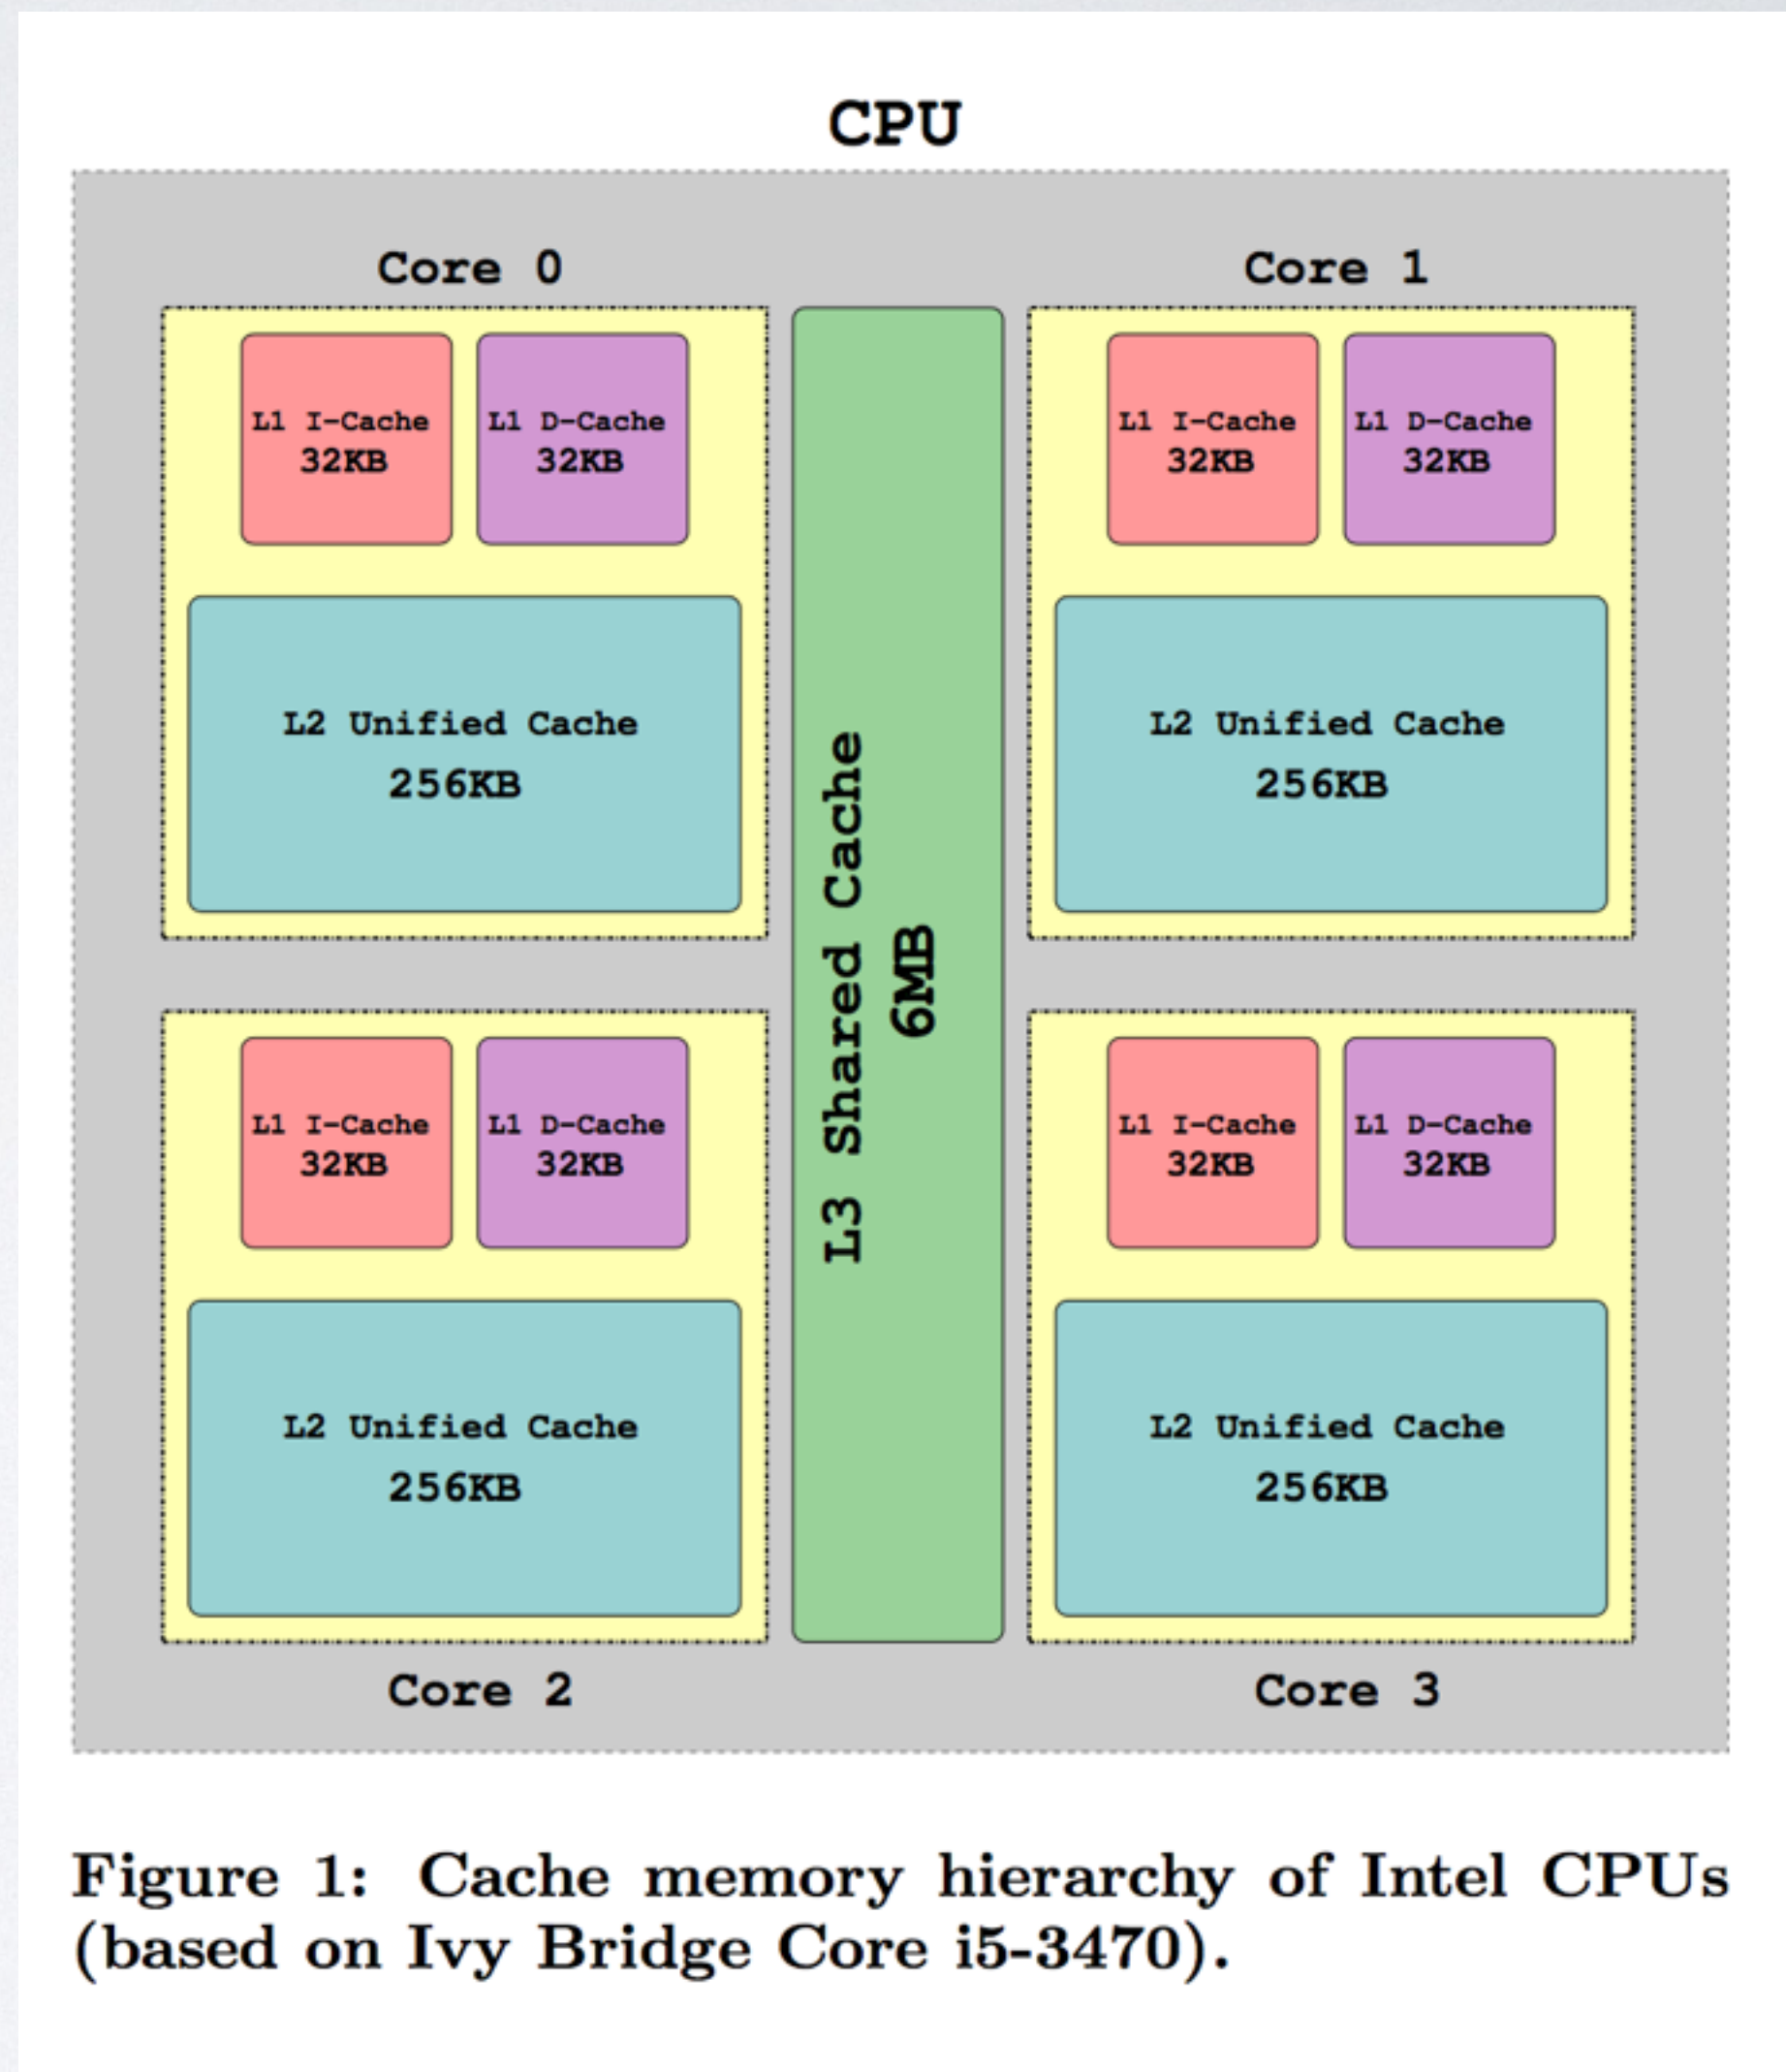


# INTRODUÇÃO



# CACHE ATTACKS

- CPUs Intel
- Cache inclusivo
- Cache L3 compartilhado entre núcleos





# PRIME + PROBE

- Eviction sets
- “Preparo e consulta”
- Deixar cache em estado conhecido
- Consultar cache após operações do usuário
- Encontrar regiões interessantes de cache

---

**Algorithm 1** Profiling a Cache Set.

---

Let  $S$  be the set of currently unmapped page-aligned addresses, and address  $x$  be an arbitrary page-aligned address in memory.

1. Repeat  $k$  times:

- (a) Iteratively access all members of  $S$ .
- (b) Measure  $t_1$ , the time it takes to access  $x$ .
- (c) Select a random page  $s$  from  $S$  and remove it.
- (d) Iteratively access all members of  $S \setminus s$ .
- (e) Measure  $t_2$ , the time it takes to access  $x$ .
- (f) If removing  $s$  caused the memory access to speed up considerably (i.e.,  $t_1 - t_2 > \text{thres}$ ), then this address is part of the same set as  $x$ . Place it back into  $S$ .
- (g) If removing  $s$  did not cause memory access to speed up considerably, then  $s$  is not part of the same set as  $x$ .


2. If  $|S| = 12$ , return  $S$ . Otherwise report failure.

---



# DETECÇÃO DA CACHE

- Verificar quais regiões de cache foram usadas através da latência



**Figure 4:** Sample memorygram collected over an idle period of 400ms. The X axis corresponds to time, while the Y axis corresponds to different cache sets. The sample shown has a temporal resolution of 250 $\mu$ s and monitors a total of 128 cache sets. The intensity of each pixel illustrates the access latency of the particular cache set, with black representing low latency and white representing a higher latency.



# DETECÇÃO DO USUÁRIO

- Atividades de rede
- Atividades do mouse




Figure 6: Memorygrams for three popular websites (Facebook, Google, Yahoo).

Classifier Output→, Ground Truth↓	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Amazon (1)	.8	-	-	-	-	-	-	.2
Baidu (2)	.2	.8	-	-	-	-	-	-
Facebook (3)	-	-	.5	-	-	.5	-	-
Google (4)	-	-	-	1	-	-	-	-
Twitter (5)	-	-	-	-	1	-	-	-
Wikipedia (6)	-	-	.2	-	-	.8	-	-
Yahoo (7)	-	-	-	-	-	-	1	-
Youtube (8)	-	-	-	-	.4	-	-	.6

Table 2: Confusion matrix for FFT-based classifier (Safari Private Browsing).



# VULNERABILIDADE COMUM

- CPUs Intel são mais de 80% dos computadores vendidos desde 2011
- Navegadores vulneráveis também representam mais de 80% dos usuários

Brand	Hi-Res. Time Support	Typed Arrays Support	Worldwide Preva- lence
Internet Explorer	v10	v11	11.77%
Safari	v8	v6	1.86%
Chrome	v20	v7	50.53%
Firefox	v15	v4	17.67%
Opera	v15	v12.1	1.2%
Total	—	—	83.03%

**Table 4: Affected desktop browsers: minimal version and prevalence [26].**



# PROBLEMAS

- Detecção correta de rede apenas 58%
- Pequeno número de *websites*
- Código JavaScript é sempre aberto
- Operações em *background* podem afetar os resultados

Our results indicate that it is possible to reliably detect mouse and network activity. The measurement rate of our network classifier did not allow us to count individual packets, but rather monitor periods of network (in)activity. Our detector was able to correctly detect 58% of these active periods, with a false positive rate of 2.86%. The mouse de-



# SOLUÇÕES À VULNERABILIDADE

- Poderia ser evitado com modificações ao JavaScript ou à arquitetura
- Possivelmente detecção pelo *runtime* do JavaScript